

Optimizing compute graphs

author

April 8, 2015

1 Enumeration

Assumption: Only care about functions f that are multivariate polynomials, e.g.,

$$f = f(X_1, X_2, \dots, X_n).$$


I am going to ignore division and minus for now. I also don't really know how to handle the coefficients, I am somehow magically assuming they don't exist (you will see later).

Definition: Compute graph \mathcal{G} is a “representation” of some function f . Each function f is “represented” by multiple compute graphs. Consider the following “recursive” definition of \mathcal{G} : A compute graph \mathcal{G} is said to be k -computation of a multivariate polynomial f , if f satisfies

$$f = g(p_1, p_2, \dots, p_k)$$

where i) g is a k -variate polynomial, and ii) p_i are polynomials in X_1, \dots, X_n that have k -computation representations by “children” graphs of \mathcal{G} . The motivation for k -computations from practical standpoints, is that the maximum number of mults/adder is k . You can also limit the number of adds but that is a simple generalization. Since I ignore division the polynomials p_i must have degree at most that of f .

The below computation graph is a 2-computation representation of f .



I am interested enumerating all possible k -computations of some multivariate polynomial f . If we can enumerate, we can pick the one with some smallest cost.


General idea: Maintain a pool of multivariate polynomials. Build larger degree polynomials by “combining” them via k -variate polynomials. We have to limit the pool by quickly eliminating polynomials that we know cannot build f . We can use the following observation to eliminate candidates.

Observation: If a polynomial p can be used to build f , then any monomial of p must divide some monomial of f .

Example: A trite one, $f = X_1^n$. In this case, any polynomial p that is a function of any X_2, \dots is eliminated.

Example: Consider $f = X_1X_2 + X_2X_3 + X_1X_3$. In this case, any polynomial p that has multiple variates, e.g. $X_2^2X_3$, is eliminated.

Example: Why this fails with minus. Very simple. We could have the following where q could be an arbitrary large degree polynomial.



Similar bad things will happen with divide.

Example: Hardamard transform. This is the Kronecker power of the matrix

$$\begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix}$$

The 2×2 transform is given by the polynomials $X_1 + X_2$ and $X_1 - X_2$. The

4×4 transform is


$$X_1 + X_2 + X_3 + X_4 \tag{1}$$

$$X_1 - X_2 + X_3 - X_4 \tag{2}$$

$$X_1 + X_2 - (X_3 + X_4) \tag{3}$$

$$X_1 - X_2 - (X_3 - X_4) \tag{4}$$

The following 1-computation graph represents (all rows of the) Kronecker



transform (or actually any Kronecker matrix product). where f_n is a row of the $2^n \times 2^n$ Hardamard transform. We simply need to adjust for the coefficients $1/-1$. This particular computation graph will achieve the smallest depth (i.e., $\log(n)$ for each row). There are of course other compute graphs, but the polynomial that represents all rows of the $2^n \times 2^n$ transform has the general form

$$f_n = \sum_{i=1}^{2^n} \prod_{j=1}^m a_{ij} X_i,$$

where the coefficients a_{ij} come from the Kronecker product. To achieve $n \log(n)$ total operations, there must be term reuse (recall the butterfly structures) so we would have to optimize 2^n compute graphs simulataneously. In this case we only need to maintain the pool of univariate polynomials in X_1, X_2, \dots, X_n .